

Model Compression with Exact Budget Constraints via Riemannian Manifolds

Michael Helcig[†]
ETH Zürich

Dan Alistarh
IST Austria

Abstract

Assigning one of K options to each of N groups under a total cost budget is a recurring problem in efficient AI, for instance in mixed-precision quantization, non-uniform pruning, and expert selection. The objective (model loss) depends on all assignments jointly and does not decompose across groups, which means that combinatorial solvers can only optimize proxy objectives. Methods such as evolutionary search evaluate the actual loss but lack gradients, while penalty-based methods enforce the budget only approximately and can require heavy hyperparameter tuning. We present a new approach to this problem by showing that under softmax relaxation, the budget constraint defines a smooth Riemannian manifold in logit space with unusually clean geometry. Specifically, the normal vector is available in closed form, shifting logits along the cost vector changes expected cost monotonically, and the vector transport reduces to a single inner product. Building on this, we propose Riemannian Constrained Optimization (RCO), which wraps tangent projection, binary-search retraction, and momentum transport around a standard Adam step. Combined with Gumbel straight-through estimation and budget-constrained dynamic programming for discrete feasibility, RCO provides first-order optimization of the actual loss under exact budget enforcement, with no constraint-related hyperparameters. RCO exceeds or matches the performance of state-of-the-art methods in both synthetic problems and realistic LLM compression settings, often at considerably lower wall-clock cost. Source code is available at <https://github.com/IST-DASLab/RCO>.

1 Introduction

Constrained optimization over discrete choices is a pervasive problem in efficient machine learning. A common instance is **budget-constrained discrete assignment**: choose one of K options for each of N target groups (e.g., a compression level for each layer of a neural network), minimizing an objective that depends on all assignments jointly (the model accuracy loss), subject to a total cost constraint (the total model size). The joint dependence between groups, e.g., the interactions between layers, makes the objective non-decomposable, so it cannot simply be written as a sum of per-group terms. Dynamic programming (DP), the standard tool for budget-constrained assignment, relies precisely on this additive structure, and can therefore only optimize proxy objectives. Concretely, variants of this problem have been used in model compression where each layer or Transformer block can be assigned a different compression degree while minimizing model loss, e.g., for mixed-precision quantization [Dong et al., 2019, 2020, Yao et al., 2021], non-uniform pruning [Frantar and Alistarh, 2022, Yin et al., 2024], MoE expert pruning [Lu et al., 2024, Lasby et al., 2025, Liu et al., 2026], and even MoE routing under load-balancing constraints [Fedus et al., 2021, Zhou et al., 2022].

Sensitivity methods address the non-decomposability by scoring groups independently, giving each layer a “sensitivity” and finding a suitable allocation via dynamic programming or integer linear

[†]Corresponding author: mhelcig@ethz.ch




Figure 1: Recovery vs. baselines across experiments: RCO matches or exceeds every baseline at lower wall-clock. Above each bar: recovery (% , bold) and wall-clock time (small).

programming [Dong et al., 2019, 2020, Yao et al., 2021, Yin et al., 2024, Frantar and Alistarh, 2022, Li et al., 2023, Malinovskii et al., 2025, Duanmu et al., 2025]. However, sensitivity-aware allocation was shown to fail at high compression rates by Sieberling et al. [2025]. Directly evaluating the actual model loss under compression combinations avoids this, but existing methods require hundreds of forward passes to converge [Sieberling et al., 2025, Liu et al., 2026]. Gradient-based search through continuous relaxation could be more efficient, but the standard approach [Cai et al., 2019, Wu et al., 2019, Huang et al., 2022] adds a quadratic penalty $\lambda (C(\alpha) - B)^2$ on the deviation between the total expected cost $C(\alpha)$ and the budget B . This never satisfies the budget exactly, and the penalty coefficient λ needs to be carefully tuned: too small and the budget is violated, too large and the constraint gradient dominates. On hard instances, as we show in Section 4.1, Lagrangian methods oscillate. Projection-free methods such as Frank-Wolfe maintain feasibility by construction [Nayman et al., 2021] but do not transport adaptive optimizer state.

We parameterize each group’s compression assignment by logits $\alpha_i \in \mathbb{R}^K$ with softmax probabilities $\mathbf{p}_i = \text{softmax}(\alpha_i)$, and let $\mathbf{c} = (c_1, \dots, c_K) \in \mathbb{R}^K$ collect the option costs and w_i the weight of group i . The total expected cost is $C(\alpha) = \sum_i w_i \langle \mathbf{p}_i, \mathbf{c} \rangle$, a smooth function of the concatenated logits $\alpha \in \mathbb{R}^{NK}$. Given a calibration loss L and a budget B , the problem is $\min_{\alpha} L(\alpha)$ subject to $C(\alpha) = B$. The logits α are the optimization variables; $\mathbf{c}$, w_i , and B are fixed inputs.

The feasible set $\mathcal{M} = \{C(\alpha) = B\}$ has well-behaved geometric structure. The gradient ∇C is everywhere nonzero when options have distinct costs, so by the regular value theorem [Boumal, 2023, Theorem 3.2] $\mathcal{M}$ is a smooth $(NK-1)$ -dimensional Riemannian submanifold of $\mathbb{R}^{NK}$. We call this the *budget manifold*; one can optimize directly on it, projecting gradients onto its tangent plane (the subspace of budget-preserving directions) and retracting (projecting back) onto its surface after each step, enforcing the budget exactly at every iterate.

The efficiently-accessible geometry of $\mathcal{M}$ rests on the softmax parameterization:

- **Closed-form normal.** The constraint gradient is $(\nabla C)_{ik} = w_i p_{ik} (c_k - \mathbb{E}_{p_i}[c])$ (Eq. 2); projecting onto the tangent plane costs a single inner product.
- **Monotonic retraction.** Shifting logits along $\mathbf{c}$ changes C monotonically with derivative $\sum_i w_i \text{Var}_{p_i}[c] > 0$ (Proposition 3); binary search therefore returns any iterate to $\mathcal{M}$ exponentially fast.
- **Cheap transport.** Since $\mathcal{M}$ has codimension one, vector transport (adjusting the optimizer’s momentum to the new tangent plane) is another inner product.

The resulting algorithm wraps these three operations around a standard Adam step, giving first-order optimization of the end-to-end loss under exact budget enforcement at negligible overhead.

Contributions. In summary, we introduce the following main contributions:

- **The budget manifold** (Section 2). The level set $\{C(\alpha) = B\}$ is a classical m -flat submanifold; under the Euclidean (rather than Fisher) ambient metric it has closed-form normals, monotonic retraction (Proposition 3), and single-inner-product transport.
- **Riemannian Constrained Optimization (RCO)** (Section 3). We wrap projection, retraction, and transport around Adam with no constraint-related hyperparameters. Gumbel-STE with budget-constrained DP handles discrete forward-pass feasibility; the manifold handles continuous backward-pass feasibility.
- **Constraint robustness under biased gradients** (Section 3, Appendix A.6). Tangent projection eliminates the constraint-normal component of any STE gradient bias (Proposition 6), so the budget is enforced exactly regardless of gradient quality.
- **Empirical validation** (Section 4). On synthetic knapsack ($N=1000$, $K=32$), RCO reaches 1% of the DP optimum where vanilla Lagrangian plateaus at 37.6%; on LLM compression, RCO matches or exceeds evolutionary search at $3\text{--}16\times$ lower wall-clock.

2 The Budget Manifold

We briefly review the relevant concepts from Riemannian geometry; for a textbook treatment, see Boumal [2023]. A *manifold* is a smooth surface that may be curved globally but locally resembles flat Euclidean space. A *Riemannian manifold* additionally carries an inner product on each tangent space, providing notions of length, angle, and gradient on the surface. At each point, the *tangent plane* is the subspace of directions tangent to the surface, and the *normal vector* points perpendicular to it. Optimizing on a manifold requires three operations: *tangent projection* (restricting gradients to the tangent plane), *retraction* (mapping an iterate that has drifted off the surface back onto it), and *vector transport* (moving vectors such as optimizer momentum from one tangent plane to another as the iterate moves along the surface).

Consider N groups, each to be assigned one of K discrete options. Group i carries weight $w_i > 0$ (e.g., the number of parameters in a network layer), and option k has cost $c_k > 0$ (e.g., bits per parameter). We parameterize the assignment distribution for group i by logits $\alpha_i \in \mathbb{R}^K$ with $\mathbf{p}_i = \text{softmax}(\alpha_i)$, and define the total expected cost as

$$C(\alpha) = \sum_{i=1}^N w_i \sum_{k=1}^K p_{ik} c_k = \sum_{i=1}^N w_i \langle \mathbf{p}_i, \mathbf{c} \rangle. \quad (1)$$

Given a budget B with $\min_k c_k < B / \sum_i w_i < \max_k c_k$, the feasible set is $\mathcal{M} = \{\alpha \in \mathbb{R}^{NK} : C(\alpha) = B\}$, and the goal is to minimize an objective $L(\alpha)$ on $\mathcal{M}$.

Proposition 1 (Manifold structure). *C is smooth. If the costs $c_1, \dots, c_K$ are not all equal, then $\nabla C(\alpha) \neq 0$ for all α , and $\mathcal{M}$ is a smooth $(NK-1)$ -dimensional submanifold of $\mathbb{R}^{NK}$.*

The proof is in Appendix A.1. The gradient is nonzero because softmax produces strictly positive probabilities, so every group has positive cost variance. Concretely, $\mathcal{M}$ is a smooth surface in logit space: at every point it has a well-defined tangent plane (budget-preserving directions) and a normal direction (the budget-changing direction), which are the two ingredients needed for constrained optimization on the surface.

Proposition 2 (Normal vector). *The gradient of C with respect to α is*

$$(\nabla C(\alpha))_{ik} = w_i p_{ik} (c_k - \mathbb{E}_{p_i}[c]), \quad (2)$$

where $\mathbb{E}_{p_i}[c] = \sum_k p_{ik} c_k$ (Appendix A.1). *This vector is normal to $\mathcal{M}$ at every point.*

Each entry measures how much option k deviates from group i ’s current expected cost, scaled by the probability and group weight.

We optimize L on $\mathcal{M}$ by wrapping three operations around a standard Adam step (Figure 2). First, we project the loss gradient $\mathbf{g} = \nabla L(\alpha)$ onto the tangent space of $\mathcal{M}$ by subtracting its component along the normal $\mathbf{n} = \nabla C(\alpha)$: $\mathbf{g}_{\text{tan}} = \mathbf{g} - (\langle \mathbf{g}, \mathbf{n} \rangle / \|\mathbf{n}\|^2) \mathbf{n}$. This removes the budget-changing direction, so the optimizer sees only budget-preserving gradients. We project before the Adam update

Input: Logits $\alpha_0 \in \mathcal{M}$, loss L , costs $\mathbf{c} \in \mathbb{R}^K$, weights $\mathbf{w} \in \mathbb{R}^N$, budget B , steps T , Gumbel samples g , temperature schedule $\{\tau_t\}$ 1: Initialize Adam state: $\mathbf{m} = \mathbf{0}$, $\mathbf{v} = \mathbf{0}$

2: **for** $t = 1, \dots, T$ **do**

 // Forward pass: Gumbel-STE + DP, averaged over g samples

 3: $\mathbf{g} \leftarrow \mathbf{0}$

 4: **for** $j = 1, \dots, g$ **do**

 5: Sample $G_{ik}^{(j)} \sim \text{Gumbel}(0, 1)$; set $\hat{\alpha}_{ik}^{(j)} = (\alpha_{ik} + G_{ik}^{(j)})/\tau_t$

 6: $\mathbf{z}^{(j)} \leftarrow \text{DP-solve}(\hat{\alpha}^{(j)}, \mathbf{c}, \mathbf{w}, B)$ Eq. (4)

 7: $\hat{\mathbf{p}}^{(j)} \leftarrow \text{softmax}(\hat{\alpha}^{(j)})$; $\mathbf{g} \leftarrow \mathbf{g} + \frac{1}{g} \nabla_{\alpha} L(\text{sg}(\mathbf{z}^{(j)} - \hat{\mathbf{p}}^{(j)}) + \hat{\mathbf{p}}^{(j)})$ STE at $\hat{\alpha}^{(j)}$; contributes $1/\tau_t$

 8: **end for**

 // Tangent projection

 9: $\mathbf{n} \leftarrow \nabla C(\alpha)$ Eq. (2)

 10: $\mathbf{g} \leftarrow \mathbf{g} - (\langle \mathbf{g}, \mathbf{n} \rangle / \|\mathbf{n}\|^2) \mathbf{n}$

 // Optimizer step + retraction

 11: $\alpha \leftarrow \text{ADAM}(\alpha, \mathbf{g}, \mathbf{m}, \mathbf{v})$

 12: Binary search for t^* s.t. $C(\alpha + t^* \tilde{\mathbf{c}}) = B$; set $\alpha_{ik} \leftarrow \alpha_{ik} + t^* c_k$ Prop. 3

 // Momentum transport

 13: $\mathbf{n}' \leftarrow \nabla C(\alpha)$

 14: $\mathbf{m} \leftarrow \mathbf{m} - (\langle \mathbf{m}, \mathbf{n}' \rangle / \|\mathbf{n}'\|^2) \mathbf{n}'$

 15: **end for**

16: **return** DP-solve(log $\mathbf{p}$, $\mathbf{c}$, $\mathbf{w}$, B) Final discrete assignment

Algorithm 1: Riemannian Constrained Optimization (RCO)

3 The Riemannian Constrained Optimization (RCO) Algorithm

The budget manifold (Section 2) provides three geometric operations for constrained optimization: tangent projection, monotonic retraction, and momentum transport. To produce discrete assignments from continuous logits, we combine these operations with Gumbel STE and DP into a complete algorithm we call *Riemannian Constrained Optimization* (RCO).

The Gumbel-STE introduces gradient bias not covered by the convergence analysis of Section 2, but the tangent projection eliminates its constraint-normal component (Proposition 6), so the budget is enforced exactly regardless of bias; the full Adam-with-annealing pipeline is validated empirically.

3.1 Algorithm

RCO enforces the budget in two complementary spaces: budget-constrained DP gives discrete feasibility in the forward pass, the manifold projection gives continuous feasibility in the backward pass. We retain the setting from Section 2: N groups, each with K options of cost c_k and group weights w_i , and a total cost budget B . Concretely, each forward pass draws Gumbel noise [Jang et al., 2017, Maddison et al., 2017] G_{ik} , forms perturbed logits $\hat{\alpha}_{ik} = (\alpha_{ik} + G_{ik})/\tau$ at temperature τ , and solves the constrained assignment

$$\mathbf{z}^* = \arg \max_{\mathbf{z} \in \{0,1\}^{NK}} \sum_{i,k} z_{ik} \hat{\alpha}_{ik} \text{ s.t. } \sum_i w_i \sum_k z_{ik} c_k \leq B, \sum_k z_{ik} = 1 \quad \forall i, \quad (4)$$

which is a multiple-choice knapsack problem solved exactly by DP in $O(NKB')$ time, where B' is the budget expressed in integer cost units (e.g., total bits for quantization, total kept experts for pruning) so the DP table has $B' + 1$ entries per group. In the backward pass, the STE [Bengio et al., 2013] replaces the non-differentiable $\arg \max$ with soft probabilities: the forward value z_{ik}^* is kept, but gradients flow through $\hat{p}_{ik} = \text{softmax}(\hat{\alpha}_i)_k$, the softmax of the same perturbed logits that produced $\mathbf{z}^*$. This ensures the surrogate concentrates on the sampled mode, so the STE bias vanishes as $\tau \rightarrow 0$ and independent Gumbel samples yield independent Jacobians; the unperturbed $\text{softmax}(\alpha_i)$ would decouple the surrogate from the sampled assignment and suppress both effects. The tangent projection (Section 2) removes the normal component from this gradient before it reaches the optimizer, eliminating any constraint-normal STE bias (Proposition 6) and preventing it from accumulating in Adam’s momentum. After optimization, the final assignment is extracted by solving (4) with $\hat{\alpha}_{ik} = \log p_{ik}$ (no noise, no temperature). Algorithm 1 gives the complete procedure.

Initialization. The algorithm assumes $\alpha_0 \in \mathcal{M}$. If the initial logits do not satisfy $C(\alpha_0) = B$ (e.g., when initialized from heuristic scores), one application of the binary-search retraction projects them onto $\mathcal{M}$ before the first step.

Temperature and variance reduction. The temperature τ in Algorithm 1 (line 5) controls exploration versus exploitation. We anneal exponentially $\tau_t = \max(\tau_{\min}, \tau_0 \cdot (\tau_{\min}/\tau_0)^{t/T})$, from $\tau_0 = 1.0$ to $\tau_{\min} = 0.01$. High temperature produces diverse DP samples and diffuse soft distributions; low temperature concentrates probability on the emerging optimum, tightening the STE approximation. The inner loop (lines 4–8) averages gradients over g independent Gumbel samples per step for variance reduction, where $g \geq 1$ is a hyperparameter (Appendix E ablates this choice).

Setup for applying RCO to LLM compression (KL objective, weight assembly via GPTQ residuals, STE backward pass) is detailed in Appendix B.

4 Experiments

4.1 Controlled validation: multiple-choice knapsack

We first isolate the manifold constraint handling from the stochastic gradient estimation. The multiple-choice knapsack problem (MCKP) admits closed-form gradients and an exact DP solution, so we can compare four constraint handling methods (manifold equality, manifold with slack variable, Lagrangian, augmented Lagrangian) on the same gradient and optimizer, varying only how they enforce the budget. We generate 13 scenarios (3 random instances per scenario, 5000 steps each) spanning correlated costs, tight budgets, adversarial instances, under-budget optima, and a large-scale instance ($N=1000$, $K=32$). Figure 3 shows convergence on the largest scenario; Table 21 (Appendix G) reports gap and violation across all scenarios.




Figure 3: MCKP, *huge* scenario ($N=1000$, $K=32$, three random instances). Left: gap to DP optimum. Center: raw constraint violation. Right: $|\text{violation}|$ on log scale; manifold ($\sim 10^{-9}$) vs. Lagrangian ($\sim 10^{-1}$). Additional scenarios in Appendix G.

Three patterns emerge. First, manifold projection satisfies the budget exactly ($|C(\alpha) - B| < 10^{-8}$) in every scenario, while Lagrangian methods maintain violations of order 10^{-1} . Second, at fixed compute the manifold’s advantage at scale is dramatic. On *huge*, the manifold reaches within 1% of the DP optimum in 594 steps; vanilla Lagrangian plateaus at a 37.6% gap from step ~ 500 onward and never improves; augmented Lagrangian narrows the gap monotonically but does not reach 1% within 5000 steps (Table 22, Appendix G). Third, when the true optimum costs less than the budget (*cheap optimal*, *mixed slack*), equality-constrained methods waste budget on inferior options; the slack variable (Section 2, Appendix A.7) recovers the DP optimum.

4.2 MoE expert pruning

We apply the manifold algorithm to MoE expert pruning, where each expert is a group with two options (keep or prune) at binary costs and the budget fixes the total number of pruned experts. EvoESAP [Liu et al., 2026] is contemporaneous work that searches over per-layer counts while holding within-layer pruning order fixed by an importance criterion (we use REAP); RCO instead

optimizes per-expert logits jointly, free to deviate from REAP ordering. We use REAP scores only to initialize the logits.

Setup. We evaluate on OLMoE-1B-7B (64 experts/layer), Qwen3-30B-A3B (128 experts/layer), and Qwen3-Coder-Next (512 experts/layer). Baselines are REAP [Lasby et al., 2025] (uniform budget allocation, REAP ordering within each layer) and EvoESAP [Liu et al., 2026] (evolutionary search over layer budgets, REAP ordering within each layer). RCO uses the same REAP scores only to initialize its per-expert logits; the search itself is free to deviate from the REAP ordering and to redistribute budget across layers. All reported pruned-model numbers are evaluated as-is; no post-pruning finetuning of routers, experts, or other parameters is performed. We report Avg, the unweighted mean of eight standard benchmarks. All EvoESAP numbers are our own reproductions under a matched protocol (same calibration data, fitness, eval harness, and compute environment as RCO); details in Appendix F.1.

Qwen3-30B-A3B. On Qwen3-30B-A3B (Figure 1; full per-benchmark breakdown in Appendix C, Table 6), RCO recovers 96% of the uncompressed baseline average at 25% pruning, vs. EvoESAP’s 90% (71.0 vs. 66.5 Avg), in ~ 85 min versus 5.2 h; at 50% pruning RCO maintains a +2.0 point Avg advantage over EvoESAP at the same wall-clock ratio.

OLMoE-1B-7B. On OLMoE-1B-7B at 25% expert sparsity, RCO at 50 steps (~ 10 min) already exceeds EvoESAP (0.597 vs. 0.581, searched for 5.8 h); at 300 steps it reaches 0.608 (+2.7 points, $5\times$ faster). Full per-benchmark iteration sweep in Appendix C (Table 5).

Qwen3-Coder-Next. Table 1 shows expert pruning on Qwen3-Coder-Next, an 80B3A MoE model with 512 experts/layer, across calibration domain and budget allocation. Nonuniform allocation is crucial at high sparsity: at 50%, it recovers 97% of HumanEval versus 55% for uniform, and at 25% both nonuniform variants match the full model. Calibration domain induces a trade-off: coding calibration preserves code generation but hurts general knowledge, while general calibration does the opposite (Appendix C.4).

Table 1: Qwen3-Coder-Next expert pruning (512 experts/layer). Uniform: fixed expert count per layer; nonuniform: RCO redistributes the budget. HE: HumanEval, MB: MBPP (pass@1); Avg: mean of eight general benchmarks. Bold: best per sparsity level. All values in %.

Sparsity	Cal.	Alloc.	Coding		General								Avg
			HE	MB	ARC-C	ARC-E	BoolQ	HSwag	MMLU	OBQA	RTE	Wino	
0%	Full model		74.4	76.4	60.6	82.1	88.5	77.5	76.7	43.0	76.5	66.6	71.4
25%	coding	uniform	68.3	68.8	50.1	72.2	86.4	69.0	71.0	38.0	72.9	65.5	65.6
	coding	nonunif.	74.4	67.8	46.2	66.2	85.1	66.5	68.0	36.2	77.6	64.2	63.8
	general	uniform	4.3	4.6	60.0	80.7	87.6	78.5	70.4	45.2	75.1	67.7	70.7
	general	nonunif.	6.1	5.8	61.8	82.2	88.2	77.6	71.2	44.2	76.2	69.9	71.4
50%	coding	uniform	40.9	53.4	40.3	64.1	78.9	57.8	56.4	35.0	67.1	61.6	57.7
	coding	nonunif.	72.0	69.0	35.6	55.5	77.6	54.8	54.3	34.0	64.6	60.3	54.6
	general	uniform	0.0	1.8	54.1	77.1	83.9	70.9	61.0	42.8	67.5	65.8	65.4
	general	nonunif.	1.2	1.0	52.6	76.2	84.2	70.8	59.5	41.4	67.5	63.5	64.4

4.3 Mixed-precision quantization

We apply RCO to mixed-precision quantization on Qwen3-8B, assigning one of seven bitwidths (2–8) to each of 252 linear layers. Layers are quantized via GPTQ [Frantar et al., 2023]; the manifold search optimizes the bitwidth assignment to minimize calibration KL divergence (Eq. 6). Calibration uses 256 FineWeb-Edu sequences (seq_len=2048). Evaluation reports perplexity on two held-out corpora (FineWeb-Edu, C4).

Baselines. Table 2 compares three RCO configurations against EvoPress [Sieberling et al., 2025] (evolutionary search over actual model loss, 100 generations), IMPQ [Zhao et al., 2025] (Shapley-based surrogate with pairwise layer interactions, solved via MILP), and dynamic HIGGS (linear

surrogate, solved via DP) [Malinovskii et al., 2025] across four average bitwidths (2.25–4.0). All EvoPress numbers for Qwen3-8B are our own reproductions under a matched protocol; Sieberling et al. [2025] do not include Qwen3-8B in their reported experiments. Details in Appendix F.2.

At high compression (2.25 bits), RCO reduces FW perplexity by 36% over HIGGS (20.47 vs. 32.07) and by 4% over EvoPress (20.47 vs. 21.40) at $\sim 10\times$ lower wall-clock cost. At 2.5 bits, RCO surpasses EvoPress (FW 15.43 vs. 15.64) at $\sim 6\times$ lower wall-clock cost (117 minutes vs. 11–14 hours). At 3.5–4.0 bits the problem is easy enough that even surrogate methods solve it well (HIGGS within 3% of RCO), and RCO’s advantage reduces to wall-clock time.

Table 2: Qwen3-8B mixed-precision quantization. Perplexity ($\downarrow$) on FineWeb-Edu (FW) and C4; FP16: FW=10.96, C4=17.20. Same GPTQ-quantized weights per layer/bitwidth. [†]Binary bitwidth; Shapley surrogate, MILP. [‡]Linear surrogate, DP.

Method	2.25 bits		2.5 bits		3.5 bits		4.0 bits		Wall
	FW	C4	FW	C4	FW	C4	FW	C4	
RCO ($g=4, T=200$)	20.60	32.44	15.78	24.73	11.41	17.86	11.17	17.52	62 m
RCO ($g=32, T=50$)	20.66	32.94	15.43	24.00	11.46	18.00	11.14	17.40	117 m
RCO ($g=16, T=200$)	20.47	32.45	15.45	23.79	11.45	17.91	11.16	17.46	215 m
EvoPress (100 gen)	21.40	33.92	15.64	24.63	11.50	18.07	11.16	17.58	11–14 h
IMPQ [†]	24.18	38.68	18.79	29.84	12.31	19.09	11.42	17.76	~ 2 h
HIGGS [‡]	32.07	53.61	20.08	31.15	11.65	18.22	11.28	17.70	~ 2 h

Ablations. Appendix E reports sweeps over nine hyperparameters. Two findings stand out. First, Gumbel sample count per step is the most important hyperparameter: increasing from $g=1$ to $g=4$ reduces FW perplexity by 0.85, with diminishing and non-monotonic gains through $g=32$. Second, sample efficiency dominates step count: $g=32$ with 50 steps matches $g=16$ with 200 steps at half the forward passes, because the STE gradient benefits more from diverse samples than from additional iterations.

4.4 Mixed-precision quantization for MoE

We compare RCO against MxMoE [Duanmu et al., 2025], a sensitivity-driven per-layer ILP, on mixed-precision weight quantization of Qwen1.5-MoE-A2.7B (24 layers, 60 routed plus one shared expert per layer) using a shared GPTQ weight database (bitwidths 2–8). Table 3 reports PPL on Wikitext-2, C4, and FineWeb-Edu and per-task zero-shot accuracy on six lm-eval-harness tasks at three expert-bit budgets, with uniform-bitwidth references. RCO and RCO* outperform MxMoE on every aggregate metric at every target despite a strictly smaller search space, and at 3.5 bit RCO matches uniform 4-bit quality, saturating against the FP16 floor; full setup and per-target discussion are in Appendix D.

Table 3: Qwen1.5-MoE-A2.7B mixed-precision quantization at 2.5 and 3.5 bit (3.0 bit block in Table 9, Appendix D). PPL ($\downarrow$) on Wikitext-2 (W2), C4, FineWeb-Edu (FW); zero-shot accuracy ($\uparrow$, decimals) on six lm-eval-harness tasks; Avg is the unweighted mean. RCO*: attention locked at FP16 (matches MxMoE). RCO: full search, attention quantized at the same total memory. Bold: best per column within each bit budget.

Bits	Method	Perplexity ($\downarrow$)				Zero-shot accuracy ($\uparrow$)						
		W2	C4	FW	Avg PPL	PIQA	HSwag	ARC-E	ARC-C	Wino	LAMBADA	Avg ZS
FP16	full model	6.79	10.05	9.07	8.64	.805	.773	.690	.445	.695	.713	.687
2.5 bit	MxMoE	7.90	12.44	10.28	10.21	.786	.746	.659	.416	.653	.662	.653
	RCO*	7.47	11.57	9.81	9.61	.793	.750	.652	.411	.651	.677	.656
	RCO	7.40	11.44	9.74	9.53	.792	.752	.667	.409	.665	.676	.660
3.5 bit	MxMoE	6.94	10.35	9.23	8.84	.801	.768	.659	.433	.686	.697	.674
	RCO*	6.90	10.25	9.18	8.77	.799	.770	.687	.445	.673	.705	.680
	RCO	6.88	10.23	9.17	8.76	.801	.771	.686	.453	.684	.706	.683
4 bit	uniform	6.91	10.24	9.20	8.78	.809	.771	.679	.440	.681	.702	.680

5 Related Work

Budget constraints in ML are predominantly enforced via penalty or Lagrangian methods. In neural architecture search, ProxylessNAS [Cai et al., 2019], FBNet [Wu et al., 2019], and SDQ [Huang et al., 2022] add expected-cost penalties to the objective; MnasNet [Tan et al., 2019] and HAQ [Wang et al., 2019] use RL reward shaping; in constrained RL, primal-dual methods update Lagrange multipliers alongside policy parameters [Achiam et al., 2017, Tessler et al., 2019, Stooke et al., 2020]. All satisfy constraints only approximately and require tuning a multiplier or schedule. RC-DARTS [Jin et al., 2020] shares the constraint form but enforces it through a decaying Lagrangian penalty, without identifying the manifold structure. HardCoRe-NAS [Nayman et al., 2021] is a notable exception, enforcing a hard latency constraint via block-coordinate Frank-Wolfe on a linearized convex polytope, but this restricts the optimizer to Frank-Wolfe updates and does not transport optimizer state across steps. Riemannian optimization restricts iterates to a smooth manifold to enforce constraints exactly [Absil et al., 2008, Boumal, 2023], with efficient algorithms for Stiefel, Grassmann, SPD, and hyperbolic manifolds [Bonnabel, 2013, Bécigneul and Ganea, 2019, Zhang and Sra, 2016, Nickel and Kiela, 2017], and for doubly stochastic matrices via Sinkhorn retraction [Douik and Hassibi, 2019]. The probability simplex and its m -flat submanifolds are classical in information geometry [Amari and Nagaoka, 2000, Lebanon, 2005], typically equipped with the Fisher metric; our Euclidean ambient metric (Section 2) is what makes the operations cheap. The closest prior work, Douik and Hassibi [2019], addresses linear constraints via iterative Sinkhorn retraction; the budget manifold instead admits scalar monotone retraction (Proposition 3). To our knowledge, this specific Riemannian manifold has not previously been used as an optimization target in the context we consider.

In non-uniform model compression, sensitivity methods [Dong et al., 2019, 2020, Yao et al., 2021, Yin et al., 2024, Li et al., 2023] score layers independently and allocate via DP or ILP, assuming the loss decomposes; HIGGS [Malinovskii et al., 2025] formalizes this via a per-layer MSE-to-perplexity linearity theorem, and IMPQ [Zhao et al., 2025] extends the surrogate with pairwise Shapley interactions. Sieberling et al. [2025] showed the decomposability assumption fails at high compression and proposed evolutionary search over actual model loss; Liu et al. [2026] (EvoESAP) extends this to MoE expert pruning. Both are zero-order and require hundreds of full model evaluations; we compare against EvoPress (Section 4.3) and EvoESAP (Section 4.2). For MoE mixed-precision quantization, MxMoE [Duanmu et al., 2025] computes per-(expert, weight-block) sensitivities and allocates via a per-layer ILP, again committing to a decomposable surrogate before observing model loss; we compare against MxMoE on Qwen1.5-MoE-A2.7B in Section 4.4. Orthogonally, Gumbel-softmax [Jang et al., 2017, Maddison et al., 2017] and the straight-through estimator [Bengio et al., 2013] enable gradient flow through discrete samples, and differentiable combinatorial solvers [Berthet et al., 2020, Vlastelica et al., 2020] extend this to structured problems; these address differentiability of discrete choices, not budget enforcement.

Prior analyses of STE bias characterize it via Taylor expansion in the categorical case [Liu et al., 2023] or zero-temperature analysis of the Gumbel-Softmax family [Shekhovtsov, 2023]; convergence guarantees for STE training are limited to two-layer networks with binary activations [Yin et al., 2019]. None of these address the projected bias relevant under manifold constraints.

6 Conclusion and Limitations

The level set of expected cost under softmax parameterization is a well-behaved Riemannian submanifold: normals are available in closed form, retraction reduces to binary search on a monotone scalar function, and vector transport is a single inner product. RCO wraps these three operations around Adam and enforces the budget exactly at every iterate, with no constraint-related hyperparameters. On synthetic knapsack, where decomposability isolates the manifold from the Gumbel-STE, RCO recovers DP-optimal solutions that Lagrangian methods miss by a wide margin. On LLM compression, where the loss is non-decomposable, RCO matches or exceeds evolutionary search at a fraction of the wall-clock cost.

Limitations. The forward-pass DP runs in $O(NKB')$ time, negligible against LLM forward passes but a bottleneck for very large option sets. STE bias makes full-algorithm convergence empirical rather than provable. Nonlinear costs (e.g., inference latency) lose retraction monotonicity, requiring iterative root-finding.

References

- P.-A. Absil, R. Mahony, and R. Sepulchre. *Optimization Algorithms on Matrix Manifolds*. Princeton University Press, 2008.
- Joshua Achiam, David Held, Abbas Tamar, and Pieter Abbeel. Constrained policy optimization. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *PMLR*, pages 22–31, 2017.
- Shun-ichi Amari and Hiroshi Nagaoka. *Methods of Information Geometry*, volume 191 of *Translations of Mathematical Monographs*. American Mathematical Society, 2000.
- Gary Bécigneul and Octavian-Eugen Ganea. Riemannian adaptive optimization methods. In *International Conference on Learning Representations*, 2019.
- Yoshua Bengio, Nicholas Léonard, and Aaron Courville. Estimating or propagating gradients through stochastic neurons for conditional computation. *arXiv preprint arXiv:1308.3432*, 2013.
- Quentin Berthet, Mathieu Blondel, Olivier Teboul, Marco Cuturi, Jean-Philippe Vert, and Francis Bach. Learning with differentiable perturbed optimizers. In *Advances in Neural Information Processing Systems*, volume 33, 2020.
- Silvère Bonnabel. Stochastic gradient descent on Riemannian manifolds. *IEEE Transactions on Automatic Control*, 58(9):2217–2229, 2013.
- Nicolas Boumal. *An Introduction to Optimization on Smooth Manifolds*. Cambridge University Press, 2023.
- Han Cai, Ligeng Zhu, and Song Han. ProxylessNAS: Direct neural architecture search on target task and hardware. In *International Conference on Learning Representations*, 2019.
- Zhen Dong, Zhewei Yao, Amir Gholami, Michael W. Mahoney, and Kurt Keutzer. HAWQ: Hessian aware quantization of neural networks with mixed-precision. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 293–302, 2019.
- Zhen Dong, Zhewei Yao, Yaohui Cai, Daiyaan Arfeen, Amir Gholami, Michael W. Mahoney, and Kurt Keutzer. HAWQ-V2: Hessian aware trace-weighted quantization of neural networks. In *Advances in Neural Information Processing Systems*, volume 33, pages 18518–18529, 2020.
- Ahmed Douik and Babak Hassibi. Manifold optimization over the set of doubly stochastic matrices: A second-order geometry. *IEEE Transactions on Signal Processing*, 67(22):5761–5774, 2019.
- Haojie Duanmu, Xiuhong Li, Zhihang Yuan, Size Zheng, Jiangfei Duan, Xingcheng Zhang, and Dahua Lin. MxMoE: Mixed-precision quantization for MoE with accuracy and performance co-design. In *Proceedings of the 42nd International Conference on Machine Learning*, PMLR, 2025.
- William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *arXiv preprint arXiv:2101.03961*, 2021.
- Elias Frantar and Dan Alistarh. SPDY: Accurate pruning with speedup guarantees. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *PMLR*, pages 6726–6743. PMLR, 2022.
- Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. GPTQ: Accurate post-training quantization for generative pre-trained transformers. In *Proceedings of the 11th International Conference on Learning Representations*, 2023.
- Hai Huang, Ao Zheng, Jianqiang Huang, Zhicheng He, and Tong He. SDQ: Stochastic differentiable quantization with mixed precision. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *PMLR*, pages 9295–9309, 2022.
- Eric Jang, Shixiang Gu, and Ben Poole. Categorical reparameterization with Gumbel-softmax. In *International Conference on Learning Representations*, 2017.

- Xiaojie Jin, Jiang Wang, Joshua Slocum, Ming-Hsuan Yang, Shengyang Dai, Shuicheng Yan, and Jiashi Feng. RC-DARTS: Resource constrained differentiable architecture search. *arXiv preprint arXiv:1912.12814*, 2020.
- Michael Lasby, Reza Bayat, Ivan Googler, Konstantinos N. Plataniotis, and Mahdi S. Hosseini. REAP: Router-weighted expert activation pruning. *arXiv preprint arXiv:2510.13999*, 2025.
- Guy Lebanon. *Riemannian Geometry and Statistical Machine Learning*. PhD thesis, Carnegie Mellon University, 2005.
- Shiyao Li, Xuefei Ning, Ke Hong, Tengxuan Liu, Luning Wang, Xiuhong Li, Kai Zhong, Guohao Dai, Huazhong Yang, and Yu Wang. LLM-MQ: Mixed-precision quantization for efficient LLM deployment. In *NeurIPS 2023 Workshop on Efficient Natural Language and Speech Processing*, pages 1–5, 2023.
- Liyuan Liu, Chengyu Dong, Xiaodong Liu, Bin Yu, and Jianfeng Gao. Bridging discrete and backpropagation: Straight-through and beyond. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Zongfang Liu, Shengkun Tang, Boyang Sun, Ping Wang, Zhiqiang Shen, and Xin Yuan. EvoESAP: Non-uniform expert pruning for sparse MoE. *arXiv preprint arXiv:2603.06003*, 2026.
- Xudong Lu, Liu Qi, Yuhui Xu, Aojun Zhou, Siyuan Huang, Bo Zhang, Junchi Yan, and Hong-Jiang Zhang. Not all experts are equal: Efficient expert pruning and skipping for mixture of experts. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.
- Chris J. Maddison, Andriy Mnih, and Yee Whye Teh. The concrete distribution: A continuous relaxation of discrete random variables. In *International Conference on Learning Representations*, 2017.
- Vladimir Malinovskii, Andrei Panferov, Ivan Ilin, Han Guo, Peter Richtárik, and Dan Alistarh. HIGGS: Pushing the limits of large language model quantization via the linearity theorem. In *Proceedings of the 2025 Conference of the North American Chapter of the Association for Computational Linguistics*, volume 1, pages 10857–10886. Association for Computational Linguistics, 2025.
- Niv Nayman, Yonathan Aflalo, Asaf Noy, and Lihi Zelnik-Manor. HardCoRe-NAS: Hard constrained differentiable neural architecture search. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *PMLR*, 2021.
- Maximilian Nickel and Douwe Kiela. Poincaré embeddings for learning hierarchical representations. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Alexander Shekhovtsov. Cold analysis of Rao-Blackwellized straight-through Gumbel-softmax gradient estimator. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *PMLR*, pages 30931–30955, 2023.
- Oliver Sieberling, Denis Kuznedelev, Eldar Kurtic, and Dan Alistarh. EvoPress: Accurate dynamic model compression via evolutionary search. In *Proceedings of the 42nd International Conference on Machine Learning*, *PMLR*, 2025.
- Adam Stooke, Joshua Achiam, and Pieter Abbeel. Responsive safety in reinforcement learning by PID Lagrangian methods. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *PMLR*, pages 9133–9143, 2020.
- Mingxing Tan, Bo Chen, Ruoming Pang, Vijay Vasudevan, Mark Sandler, Andrew Howard, and Quoc V. Le. MnasNet: Platform-aware neural architecture search for mobile. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 2820–2828, 2019.
- Chen Tessler, Daniel J. Mankowitz, and Shie Mannor. Reward constrained policy optimization. In *International Conference on Learning Representations*, 2019.
- Marin Vlastelica, Anselm Paulus, Vit Musil, Georg Martius, and Michal Rolínek. Differentiation of blackbox combinatorial solvers. In *International Conference on Learning Representations*, 2020.

- Martin J. Wainwright and Michael I. Jordan. Graphical models, exponential families, and variational inference. *Foundations and Trends in Machine Learning*, 1(1–2):1–305, 2008.
- Kuan Wang, Zhijian Liu, Yujun Lin, Ji Lin, and Song Han. HAQ: Hardware-aware automated quantization with mixed precision. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 8612–8620, 2019.
- Bichen Wu, Xiaoliang Dai, Peizhao Zhang, Yanghan Wang, Fei Sun, Yiming Wu, Yuandong Tian, Peter Vajda, Yangqing Jia, and Kurt Keutzer. FBNet: Hardware-aware efficient ConvNet design via differentiable neural architecture search. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10734–10742, 2019.
- Zhewei Yao, Zhen Dong, Zhangcheng Zheng, Amir Gholami, Jiali Yu, Eric Tan, Leyuan Wang, Qijing Huang, Yida Wang, Michael W. Mahoney, and Kurt Keutzer. HAWQ-V3: Dyadic neural network quantization. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *PMLR*, pages 11875–11886. PMLR, 2021.
- Lu Yin, You Wu, Zhenyu Zhang, Cheng-Yu Hsieh, Yaqing Wang, Yiling Jia, Gen Li, Ajay Jaiswal, Mykola Pechenizkiy, Yi Liang, Michael Bendersky, Zhangyang Wang, and Shiwei Liu. OWL: Outlier weighed layerwise sparsity for accelerating large language models. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *PMLR*, pages 57101–57115. PMLR, 2024.
- Penghang Yin, Jiancheng Lyu, Shuai Zhang, Stanley Osher, Yingyong Qi, and Jack Xin. Understanding straight-through estimator in training activation quantized neural nets. In *International Conference on Learning Representations*, 2019.
- Hongyi Zhang and Suvrit Sra. First-order methods for geodesically convex optimization. In *Conference on Learning Theory*, volume 49 of *PMLR*, pages 1617–1638, 2016.
- Junchen Zhao, Ali Derakhshan, Dushyant Bharadwaj, Jayden Kana Hyman, Junhao Dong, Sangeetha Abdu Jyothi, and Ian Harris. IMPQ: Interaction-aware layerwise mixed precision quantization for LLMs. *arXiv preprint arXiv:2509.15455*, 2025.
- Yanqi Zhou, Tao Lei, Hanxiao Liu, Nan Du, Yanping Huang, Vincent Zhao, Andrew Dai, Zhifeng Chen, Quoc Le, and James Laudon. Mixture-of-experts with expert choice routing. In *Advances in Neural Information Processing Systems*, volume 35, 2022.

A Proofs and Theoretical Details

This appendix provides formal statements and proofs for results referenced in Section 2, verifies that the algorithm components satisfy standard Riemannian optimization axioms, and derives the extensions to inequality and multiple constraints.

A.1 Manifold structure and normal vector

Proof of Proposition 1. The expected cost $C(\alpha) = \sum_i w_i \langle \text{softmax}(\alpha_i), \mathbf{c} \rangle$ is a composition of the smooth softmax with a linear function, hence smooth. By Proposition 2, the (i, k) entry of $\nabla C(\alpha)$ is $w_i p_{ik}(c_k - \mathbb{E}_{p_i}[c])$. Since softmax produces strictly positive probabilities ($p_{ik} > 0$ for all i, k),

$$\|\nabla C\|^2 = \sum_{i=1}^N w_i^2 \sum_{k=1}^K p_{ik}^2 (c_k - \mathbb{E}_{p_i}[c])^2.$$

If the costs are not all equal, then for every group i , $\mathbb{E}_{p_i}[c]$ is a strict convex combination of distinct values, so at least one $c_k \neq \mathbb{E}_{p_i}[c]$, and $p_{ik} > 0$ gives $\|\nabla C\|^2 > 0$. The regular value theorem [Boumal, 2023, Theorem 3.2] then implies that $\mathcal{M}$ is a smooth $(NK-1)$ -dimensional submanifold. $\square$

Proof of Proposition 2. The softmax Jacobian is $\partial p_{ij} / \partial \alpha_{ik} = p_{ij}(\delta_{jk} - p_{ik})$. Differentiating (1):

$$\begin{aligned} \frac{\partial C}{\partial \alpha_{ik}} &= w_i \sum_{j=1}^K c_j p_{ij} (\delta_{jk} - p_{ik}) \\ &= w_i (c_k p_{ik} - p_{ik} \sum_j c_j p_{ij}) \\ &= w_i p_{ik} (c_k - \mathbb{E}_{p_i}[c]). \end{aligned} \quad \square$$

A.2 Monotonic retraction

Proof of Proposition 3. Let $\tilde{c}_{ik} = c_k$. By the chain rule and Proposition 2:

$$\begin{aligned} \frac{d}{dt} C(\alpha + t\tilde{\mathbf{c}}) &= \sum_{i,k} w_i p_{ik}(t) (c_k - \mathbb{E}_{p_i(t)}[c]) c_k \\ &= \sum_{i=1}^N w_i \left(\sum_k p_{ik}(t) c_k^2 - (\mathbb{E}_{p_i(t)}[c])^2 \right) \\ &= \sum_{i=1}^N w_i \text{Var}_{p_i(t)}[c]. \end{aligned}$$

Since $w_i > 0$ and $p_{ik}(t) > 0$ for all k , $\text{Var}_{p_i(t)}[c] > 0$ whenever costs are not all equal. Moreover, as $t \rightarrow +\infty$ the softmax concentrates on $\arg \max_k c_k$ giving $C \rightarrow \sum_i w_i \max_k c_k$, and as $t \rightarrow -\infty$, $C \rightarrow \sum_i w_i \min_k c_k$. Strict monotonicity and these limiting values ensure that $C(\alpha + t\tilde{\mathbf{c}}) = B$ has a unique solution for any B in the feasible range, found by binary search with exponential convergence in the number of iterations. $\square$

Empirical retraction cost. With shift bracket $t \in [-50, 50]$ (range $R = 100$) and tolerance $\varepsilon = 10^{-8}$ on the cost residual, the binary search converges in $\lceil \log_2(R/\varepsilon) \rceil = 34$ iterations in t -space. The cost residual is enforced in expected-cost space, where the local slope $\sum_i w_i \text{Var}_{p_i}[c]$ varies across iterates, so the empirical iteration count differs slightly from this theoretical bound. Figure 4 reports the distribution across five MCKP scenarios (huge, large, medium, correlated tight, adversarial; 1.8k–5.0k retraction calls each). Per-call mean ranges from 29.4 on *adversarial* to 40.8 on *huge*; the maximum across all scenarios and calls is 45.




Figure 4: Binary-search iterations per retraction across five MCKP scenarios (boxes: IQR; whiskers: 1.5 IQR; points: outliers). Dashed green line: theoretical bound $\lceil \log_2(R/\varepsilon) \rceil = 34$ for the configuration used in our experiments ($R = 100$, $\varepsilon = 10^{-8}$). Dotted gray line: prior loose estimate of 50. All scenarios concentrate near the theoretical bound and stay below 45 in the worst case.

A.3 Verification of retraction axioms

A retraction on a manifold $\mathcal{M}$ is a smooth map $R: T\mathcal{M} \rightarrow \mathcal{M}$ satisfying two axioms [Absil et al., 2008, Definition 4.1]: (i) $R_\alpha(\mathbf{0}) = \alpha$ (centering), and (ii) $\frac{d}{dt} R_\alpha(t\xi)|_{t=0} = \xi$ for all $\xi \in T_\alpha\mathcal{M}$ (local rigidity).

Our retraction takes a tangent vector $\xi \in T_\alpha\mathcal{M}$, forms $\alpha' = \alpha + \xi$, and finds t^* such that $C(\alpha' + t^*\tilde{\mathbf{c}}) = B$. Smoothness of t^* as a function of ξ follows from the implicit function theorem applied to

$$F(\xi, t) = C(\alpha + \xi + t\tilde{\mathbf{c}}) - B.$$

The partial derivative $\partial F / \partial t = \sum_i w_i \text{Var}_{p_i}[c] > 0$ (Proposition 3) is nonzero everywhere, so $t^*(\xi)$ is smooth in a neighborhood of any point on $\mathcal{M}$.

Proposition 4 (Retraction axioms). *The map $R_\alpha(\xi) = \alpha + \xi + t^*(\xi)\tilde{\mathbf{c}}$, where t^* solves $C(\alpha + \xi + t^*\tilde{\mathbf{c}}) = B$, satisfies both retraction axioms.*

Proof. Centering. When $\xi = \mathbf{0}$, $C(\alpha) = B$ already holds, so $t^* = 0$ and $R_\alpha(\mathbf{0}) = \alpha$.

Local rigidity. Write $R_\alpha(t\xi) = \alpha + t\xi + t^*(t\xi)\tilde{\mathbf{c}}$. Differentiating the constraint

$$C(\alpha + t\xi + t^*(t\xi)\tilde{\mathbf{c}}) = B$$

at $t = 0$ gives

$$\langle \nabla C(\alpha), \xi \rangle + \frac{dt^*}{dt} \Big|_{t=0} \langle \nabla C(\alpha), \tilde{\mathbf{c}} \rangle = 0.$$

The first term vanishes because $\xi \in T_\alpha\mathcal{M}$, and $\langle \nabla C, \tilde{\mathbf{c}} \rangle \neq 0$ by Proposition 3, so $\frac{dt^*}{dt} \Big|_{t=0} = 0$. Therefore $\frac{d}{dt} \Big|_{t=0} R_\alpha(t\xi) = \xi$. $\square$

A.4 Riemannian gradient

The Riemannian gradient of a function $f: \mathcal{M} \rightarrow \mathbb{R}$ at $\alpha \in \mathcal{M}$ is the unique tangent vector $\text{grad } f(\alpha) \in T_\alpha\mathcal{M}$ satisfying $\langle \text{grad } f, \xi \rangle = Df(\alpha)[\xi]$ for all $\xi \in T_\alpha\mathcal{M}$ [Boumal, 2023, Proposition 3.61]. We verify that the tangent projection used in Algorithm 1 recovers it.

Proposition 5 (Projected gradient is the Riemannian gradient). *Let $L: \mathbb{R}^{NK} \rightarrow \mathbb{R}$ be smooth and let $\mathbf{n} = \nabla C(\alpha)$. Then*

$$\text{grad } L|_{\mathcal{M}}(\alpha) = \nabla L(\alpha) - \frac{\langle \nabla L(\alpha), \mathbf{n} \rangle}{\|\mathbf{n}\|^2} \mathbf{n}.$$

Proof. The tangent space is $T_\alpha \mathcal{M} = \{\xi : \langle \mathbf{n}, \xi \rangle = 0\}$. Since $\mathcal{M}$ inherits the Euclidean metric, the Riemannian gradient equals the orthogonal projection of ∇L onto $T_\alpha \mathcal{M}$. The normal space is $\text{span}(\mathbf{n})$, giving the stated formula. $\square$

This confirms that every projected-gradient step in Algorithm 1 moves in the steepest-descent direction on $\mathcal{M}$. Note that this differs from projected gradient descent (PGD), which projects the *iterate* onto a convex constraint set after each step. Here we project the *gradient* onto the tangent space and use a separate retraction to bring the iterate back to $\mathcal{M}$. This is necessary because $\mathcal{M}$ is not convex, so projecting iterates onto it is not well-defined in general.

A.5 Vector transport by projection

Exact parallel transport on implicitly defined submanifolds is expensive to compute. We instead use *vector transport by projection* [Absil et al., 2008, Section 8.1.3]: given $\xi \in T_\alpha \mathcal{M}$ and a retraction step to α' , the transported vector is

$$\mathcal{T}_{\alpha \rightarrow \alpha'}(\xi) = \xi - \frac{\langle \xi, \mathbf{n}' \rangle}{\|\mathbf{n}'\|^2} \mathbf{n}',$$

where $\mathbf{n}' = \nabla C(\alpha')$. This satisfies the vector transport axioms [Absil et al., 2008, Definition 8.1.1] and coincides with parallel transport to first order in the step size.

In Algorithm 1, we apply this transport to Adam’s first moment $\mathbf{m}$. We only transport the first moment; the second moment is a scalar scaling factor per coordinate and does not need transport.

A.6 Robustness to gradient bias

The convergence analysis of Riemannian gradient descent (Section 2) assumes exact gradients. In practice, Algorithm 1 uses biased gradient estimates from the straight-through estimator. The following result shows that the tangent projection removes the normal component of any bias before it reaches the optimizer, eliminating the first-order constraint violation that unprojected bias would cause; the remaining per-step violation is second-order, scaling with manifold curvature times $\|P_T(\tilde{\mathbf{g}})\|^2$ (which still depends on the tangential bias) and is corrected exactly by the retraction at every step.

Proposition 6 (Constraint robustness under biased gradients). *Let $\tilde{\mathbf{g}} = \nabla L(\alpha) + \mathbf{b}$ be a biased gradient estimate at $\alpha \in \mathcal{M}$, where $\mathbf{b}$ is an arbitrary bias vector. Let $\mathbf{n} = \nabla C(\alpha)$ denote the constraint normal and write $P_T(\mathbf{v}) = \mathbf{v} - \langle \mathbf{v}, \mathbf{n} \rangle / \|\mathbf{n}\|^2 \mathbf{n}$ for the tangent projection.*

- (i) **Normal bias elimination.** *The tangent projection eliminates the constraint-normal component of the bias:*

$$P_T(\tilde{\mathbf{g}}) = P_T(\nabla L) + \mathbf{b}_T,$$

where $\mathbf{b}_T = \mathbf{b} - \frac{\langle \mathbf{b}, \mathbf{n} \rangle}{\|\mathbf{n}\|^2} \mathbf{n}$ is the tangential component of the bias. The normal component $\mathbf{b}_n = \mathbf{b} - \mathbf{b}_T$ does not enter the optimizer.

- (ii) **Constraint violation bound (gradient descent).** *For a gradient step $\alpha' = \alpha - \eta P_T(\tilde{\mathbf{g}})$, the per-step constraint violation satisfies*

$$|C(\alpha') - B| \leq \frac{\eta^2}{2} \|\nabla^2 C(\alpha)\|_{\text{op}} \|P_T(\tilde{\mathbf{g}})\|^2 + O(\eta^3).$$

This bound is second-order in η and independent of $\|\mathbf{b}_n\|$. Without projection, the violation is first-order: $|C(\alpha - \eta \tilde{\mathbf{g}}) - B| = \eta |\langle \mathbf{n}, \tilde{\mathbf{g}} \rangle| + O(\eta^2)$.

- (iii) **Momentum insulation (Adam).** *In Algorithm 1, after tangent projection (line 10) and vector transport (line 14), Adam’s first moment satisfies $\langle \mathbf{m}_t, \mathbf{n}_t \rangle = 0$ for all $t \geq 0$, regardless of the bias in the gradient estimates.*

Proof. Part (i). P_T is linear (it is an orthogonal projection onto the hyperplane $\{\mathbf{v} : \langle \mathbf{v}, \mathbf{n} \rangle = 0\}$), so $P_T(\tilde{\mathbf{g}}) = P_T(\nabla L) + P_T(\mathbf{b})$. Since $P_T(\mathbf{b}) = \mathbf{b} - \frac{\langle \mathbf{b}, \mathbf{n} \rangle}{\|\mathbf{n}\|^2} \mathbf{n} = \mathbf{b}_T$, the normal component $\mathbf{b}_n$ is removed.

Part (ii). Let $\mathbf{d} = P_T(\tilde{\mathbf{g}})$. Taylor-expanding C at $\alpha \in \mathcal{M}$:

$$C(\alpha - \eta \mathbf{d}) = \underbrace{C(\alpha)}_{=B} - \eta \underbrace{\langle \nabla C(\alpha), \mathbf{d} \rangle}_{=\langle \mathbf{n}, \mathbf{d} \rangle = 0} + \frac{\eta^2}{2} \mathbf{d}^\top \nabla^2 C(\alpha) \mathbf{d} + O(\eta^3).$$

The first-order term vanishes because $\mathbf{d} \in T_\alpha \mathcal{M}$. The Hessian $\nabla^2 C$ exists by smoothness of C (Proposition 1); its operator norm is bounded because every entry is a polynomial in the softmax probabilities $p_{ik} \in [0, 1]$ with coefficients depending only on w_i and c_k , giving the stated bound. Without projection, $\mathbf{d} = \tilde{\mathbf{g}}$ and $\langle \mathbf{n}, \tilde{\mathbf{g}} \rangle \neq 0$ in general, so the violation is $O(\eta)$.

Part (iii). By induction on t . At each step, the momentum passes through two stages: the Adam update (which mixes old momentum with the new projected gradient) and the vector transport (which re-projects onto the tangent space at the retracted point). We write $\mathbf{m}_t^+$ for the momentum after the Adam update and $\mathbf{m}_t^\mathcal{T}$ for the momentum after transport.

Base case. $\mathbf{m}_0^\mathcal{T} = \mathbf{0}$, so $\langle \mathbf{m}_0^\mathcal{T}, \mathbf{n} \rangle = 0$ for any $\mathbf{n}$.

Inductive step. Suppose $\langle \mathbf{m}_{t-1}^\mathcal{T}, \mathbf{n}_t \rangle = 0$, where $\mathbf{n}_t = \nabla C(\alpha_t)$ is the normal at the current iterate. The projected gradient satisfies $\langle P_T(\tilde{\mathbf{g}}_t), \mathbf{n}_t \rangle = 0$ by construction. Adam's momentum update is a convex combination of these two orthogonal-to- $\mathbf{n}_t$ vectors: $\mathbf{m}_t^+ = \beta_1 \mathbf{m}_{t-1}^\mathcal{T} + (1 - \beta_1) P_T(\tilde{\mathbf{g}}_t)$, so $\langle \mathbf{m}_t^+, \mathbf{n}_t \rangle = \beta_1 \cdot 0 + (1 - \beta_1) \cdot 0 = 0$. After the Adam step and retraction to $\alpha'_t \in \mathcal{M}$, vector transport (Algorithm 1, line 14) projects $\mathbf{m}_t^+$ onto the new tangent space: $\mathbf{m}_t^\mathcal{T} = \mathbf{m}_t^+ - (\langle \mathbf{m}_t^+, \mathbf{n}'_t \rangle / \|\mathbf{n}'_t\|^2) \mathbf{n}'_t$, which satisfies $\langle \mathbf{m}_t^\mathcal{T}, \mathbf{n}'_t \rangle = 0$ by construction of the projection. Since $\alpha_{t+1} = \alpha'_t$, we have $\mathbf{n}_{t+1} = \mathbf{n}'_t$, completing the induction. $\square$

Remark 7 (Adam's per-coordinate scaling). Part (ii) applies to gradient descent. Adam's step direction is $\mathbf{d} = \hat{\mathbf{m}} \odot (\sqrt{\hat{\mathbf{v}}} + \epsilon)$ (element-wise), and the per-coordinate scaling does not preserve the tangent plane: $\langle \mathbf{n}, \mathbf{d} \rangle \neq 0$ even when $\langle \mathbf{m}, \mathbf{n} \rangle = 0$. To quantify how much normal leakage Adam introduces, write $D = \text{diag}(1/(\sqrt{\hat{v}_{ik}} + \epsilon))$ for the scaling matrix and $\bar{d}$ for its mean diagonal entry, and decompose D into its uniform part $\bar{d}I$ (which preserves the tangent plane) and the non-uniform residual $D - \bar{d}I$:

$$\langle \mathbf{n}, D\mathbf{m} \rangle = \mathbf{m}^\top (D - \bar{d}I) \mathbf{n} + \bar{d} \underbrace{\langle \mathbf{m}, \mathbf{n} \rangle}_{=0 \text{ by (iii)}}.$$

The projection eliminates the direct bias-driven normal component $\bar{d} \langle \mathbf{m}, \mathbf{n} \rangle$. The remaining normal leakage $|\mathbf{m}^\top (D - \bar{d}I) \mathbf{n}| \leq \|\mathbf{m}\| \|\mathbf{n}\| \|D - \bar{d}I\|_{\text{op}}$ is proportional to the non-uniformity of Adam's second-moment estimates: if D were a scalar matrix (uniform scaling), this term would vanish regardless of $\|\mathbf{m}\|$. The factor $\|\mathbf{m}\|$ still reflects the accumulated (tangential) gradient history, including tangential bias; what the projection guarantees is that bias enters this bound only through the non-uniformity factor $\|D - \bar{d}I\|_{\text{op}}$, not directly. The retraction (Proposition 3) corrects this residual exactly.

Proposition 6 establishes that the tangent projection and vector transport isolate constraint satisfaction from gradient estimation quality: the constraint-normal component of any bias is eliminated before it reaches the optimizer. The tangential component $\mathbf{b}_T$ does not cause constraint violation, but it does affect optimization quality.

A.7 Inequality constraints via slack variables

For the inequality constraint $C(\alpha) \leq B$, we introduce a scalar slack variable $s \in \mathbb{R}$ and define the augmented equality constraint

$$\hat{C}(\alpha, s) = C(\alpha) + s^2 = B. \quad (5)$$

This defines a smooth manifold $\hat{\mathcal{M}} \subset \mathbb{R}^{NK+1}$. The gradient of $\hat{C}$ in the augmented space is

$$\nabla \hat{C}(\alpha, s) = (\nabla_\alpha C(\alpha), 2s).$$

When $s > 0$, the normal has a nonzero component in the s -direction, so the tangent space includes directions that change α freely (with compensating changes in s). The optimizer can decrease cost below budget by increasing s . As $s \rightarrow 0$, the algorithm reduces to the equality-constrained case.

The tangent projection in the augmented space is

$$\text{proj}(\mathbf{g}, 0) = (\mathbf{g}, 0) - \frac{\langle \mathbf{g}, \nabla_{\alpha} C \rangle}{\|\nabla_{\alpha} C\|^2 + 4s^2} (\nabla_{\alpha} C, 2s),$$

where $\mathbf{g} = \nabla_{\alpha} L$ and the s -component of the objective gradient is zero (the objective does not depend on s).

Retraction proceeds by first updating (α, s) via the projected Adam step, then adjusting: if $C(\alpha) > B$, retract α via binary search to $C(\alpha) = B$ and set $s = 0$; otherwise set $s = \sqrt{B - C(\alpha)}$. This maintains $\hat{C}(\alpha, s) = B$ exactly.

Smoothness at $s = 0$. The retraction is continuous but not smooth at the boundary $s = 0$, because the square-root branch $s' = \sqrt{B - C(\alpha')}$ has an infinite derivative as $s' \rightarrow 0^+$. In practice, the optimizer either settles to a point with $s > 0$ (the optimum is strictly under budget) or converges to the equality-constrained manifold ($s = 0$), spending at most a few iterations near the boundary. We observe no numerical issues from this non-smoothness in any of our experiments.

A.8 Multiple equality constraints

Given q constraints $C_j(\alpha) = b_j$ for $j = 1, \dots, q$, the feasible set is $\mathcal{M} = \bigcap_j C_j^{-1}(b_j)$. Let $\mathbf{n}_j = \nabla C_j(\alpha)$ and $\mathbf{N} = [\mathbf{n}_1 \dots \mathbf{n}_q]$. If the $\mathbf{n}_j$ are linearly independent, $\mathcal{M}$ is a smooth $(NK - q)$ -dimensional submanifold with tangent space

$$T_{\alpha} \mathcal{M} = \{\xi : \mathbf{N}^{\top} \xi = \mathbf{0}\}.$$

The tangent projection becomes

$$\mathbf{g}_{\text{tan}} = \mathbf{g} - \mathbf{N}(\mathbf{N}^{\top} \mathbf{N})^{-1} \mathbf{N}^{\top} \mathbf{g},$$

which requires solving a $q \times q$ system, cheap when $q \ll NK$.

Retraction generalizes via Newton root-finding on the q -dimensional system: find $\mathbf{t} \in \mathbb{R}^q$ such that $C_j(\alpha + \sum_l t_l \tilde{\mathbf{c}}_l) = b_j$ for all j . The Jacobian entry is

$$J_{jl} = \sum_{i=1}^N w_i \text{Cov}_{p_i}[c_j, c_l],$$

available in closed form from the current softmax distributions. For $q = 1$ this reduces to the scalar $\sum_i w_i \text{Var}_{p_i}[c]$ from Proposition 3, and binary search on the same monotone function is the simpler alternative used in Section 2. Quadratic convergence typically requires 3–4 iterations for general q . When the constraints are separable (each depends on a disjoint subset of groups), independent scalar retraction per constraint suffices.

Figure 5 validates this extension on a synthetic problem with $q=16$ simultaneous resource constraints, $N=500$ groups, and $K=32$ options. The budget manifold satisfies all 16 constraints to $\sim 10^{-13}$ (the limit of double-precision arithmetic) throughout optimization. Lagrangian penalty cannot balance 16 independent multipliers and diverges to violations of order 10; augmented Lagrangian reduces this to $\sim 10^{-2}$ but still exceeds the budget manifold by ten orders of magnitude.

B Application to Non-uniform LLM Compression

We apply RCO to non-uniform LLM compression, where each network layer (or expert) is a group and the options are compression levels (bitwidths, sparsity rates, or keep/prune).

Objective. The loss $L(\mathbf{z}^*)$ in Algorithm 1 is the KL divergence between the full-precision model and the model under assignment $\mathbf{z}^*$, evaluated on a calibration set $\mathcal{D}$:

$$L(\mathbf{z}^*) = \frac{1}{|\mathcal{D}|} \sum_{x \in \mathcal{D}} \text{KL}(p_{\text{ref}}(\cdot | x) \parallel p_{\mathbf{z}^*}(\cdot | x)), \quad (6)$$

where p_{ref} is the full-precision model’s next-token distribution and $p_{\mathbf{z}^*}$ is the distribution under configuration $\mathbf{z}^*$. Reference log-probabilities are computed once and cached. Because L is non-decomposable (Section 1), minimizing it directly is what distinguishes this approach from sensitivity methods that optimize per-group proxy scores.




Figure 5: Per-constraint violation trajectories with $q=16$ simultaneous budget constraints ($N=500$, $K=32$). Each trace is one constraint; y-axis scales differ between panels. Budget manifold enforces all 16 constraints to $\sim 10^{-13}$ throughout optimization. Lagrangian methods lack the geometric structure to balance many constraints and exhibit persistent violations orders of magnitude larger.

From assignments to weights. Each forward pass evaluates the model under $\mathbf{z}^*$ from Eq. (4): for MoE expert pruning, a per-expert (keep, prune) mask at costs $(0, 1)$ that zeros pruned outputs and leaves router logits unchanged; for mixed-precision quantization, a one-hot bitwidth selector that assembles $\mathbf{W}_i = \mathbf{W}_i^{\text{ref}} + \sum_k z_{ik}^* \Delta_{ik}$ from pre-computed GPTQ residuals $\Delta_{ik} = \mathbf{W}_i^{(k)} - \mathbf{W}_i^{\text{ref}}$ [Frantar et al., 2023]. DP costs and budget are taken in integer units (bitwidths scaled by parameter fraction for quantization, binary keep/prune for expert pruning). Algorithm 1 is agnostic to grouping and to the initial logits α_0 : any partition of compressible units, any per-group option set, and any prior on α_0 leave the iteration unchanged, with feasibility absorbed by the retraction-based initialization step (Section 3). Sensitivity scores from prior work, such as REAP [Lasby et al., 2025] for expert pruning or layer-output L2 sensitivities for quantization, serve as valid initializations alongside uniform $\alpha_0 = \mathbf{0}$. The STE replaces z_{ik}^* with $\hat{p}_{ik} = \text{softmax}(\hat{\alpha}_i)_k$ in the backward pass, so gradients flow through the softmax Jacobian at the perturbed logits.

C Full Expert Pruning Results

Gradient signal under sparse routing. In MoE models, only the top- k routed experts are computed per token; the STE pruning mask is applied after routing by scaling each selected expert’s contribution. An expert’s pruning logit α_i therefore receives gradient only from tokens where expert i was routed, which is a small fraction of the calibration set. This makes the per-expert gradient signal inherently sparser than in the quantization setting, where every layer processes every token. The sparsity explains why the Gumbel sample count g is the most important hyperparameter in the expert pruning ablations: more samples per step expose each expert to more diverse pruning contexts, compensating for the low per-token coverage. Because router weights are frozen, pruned experts are masked out at inference in the same way as during optimization (the remaining top- k experts are upweighted); there is no approximation gap between the STE training objective and deployment.

C.1 OLMoE-1B-7B: sparsity sweep

Table 4 reports all eight benchmarks across sparsity levels from 5% to 50% (RCO, 300 steps).

Table 4: OLMoE-1B-7B per-benchmark results across sparsity levels (RCO, 300 steps). Avg is the unweighted mean of all eight benchmarks.

	Full	5%	10%	15%	20%	25%	50%
ARC-C	.493	.497	.489	.497	.470	.458	.334
ARC-E	.758	.760	.741	.738	.718	.698	.525
BoolQ	.768	.757	.750	.717	.707	.713	.602
HellaSwag	.806	.797	.779	.759	.735	.711	.469
MMLU	.534	.531	.526	.506	.495	.479	.301
OBQA	.468	.470	.466	.462	.450	.456	.326
RTE	.715	.726	.693	.704	.718	.755	.570
WinoGrande	.680	.690	.650	.645	.624	.596	.533
Avg	.653	.653	.637	.629	.615	.608	.458

At 5% sparsity, every individual benchmark remains within 1–2 percentage points of the full model; several (ARC-C, ARC-E, OBQA, RTE) are within noise. Degradation is monotonic for most benchmarks, with the exception of RTE, which improves under moderate pruning (0.715 $\rightarrow$ 0.755 at 25%). The largest per-benchmark drops occur at 50% sparsity on HellaSwag (−33.7 points) and MMLU (−23.3 points), consistent with these benchmarks’ sensitivity to knowledge capacity.

C.2 OLMoE-1B-7B: full iteration sweep

Table 5 reports the full per-benchmark iteration sweep on OLMoE-1B-7B at 25% sparsity.

Table 5: OLMoE-1B-7B at 25% sparsity: full per-benchmark iteration sweep. All eight benchmarks contributing to Avg.

	ARC-C	ARC-E	BoolQ	HSwag	MMLU	OBQA	RTE	Wino	Avg	Time
Full model	.493	.758	.768	.806	.534	.468	.715	.680	.653	-
REAP	.408	.676	.683	.660	.452	.406	.664	.580	.566	<1 m
EvoESAP	.427	.641	.673	.685	.474	.386	.614	.649	.581	5.8 h
RCO ($T=50$)	.442	.713	.693	.691	.476	.440	.718	.602	.597	~10 m
RCO ($T=100$)	.446	.716	.692	.713	.465	.438	.722	.608	.600	~23 m
RCO ($T=150$)	.454	.727	.697	.703	.427	.452	.668	.616	.593	~32 m
RCO ($T=200$)	.459	.701	.701	.708	.475	.420	.697	.608	.596	~44 m
RCO ($T=300$)	.458	.698	.713	.711	.479	.456	.755	.596	.608	~66 m
RCO ($T=500$)	.456	.699	.698	.711	.479	.444	.722	.612	.603	~108 m

RCO reaches 0.597 Avg at 50 steps and improves to 0.608 at 300 steps. Returns diminish beyond 300 steps (500 steps: 0.603). Individual benchmarks show non-monotonic behavior during optimization (e.g., MMLU dips at $T=150$), reflecting the stochastic nature of Gumbel-STE sampling, but the aggregate Avg trends upward through 300 steps.

C.3 Qwen3-30B-A3B: full per-benchmark comparison

Table 6 reports the full per-benchmark breakdown of RCO vs. EvoESAP on Qwen3-30B-A3B at 25% and 50% sparsity (summarized in Figure 1).

Table 6: Qwen3-30B-A3B per-expert pruning: full benchmark breakdown. RCO wins on all 8 benchmarks at 25% and on 6/8 at 50%. Largest gains at 25%: HellaSwag (+8.7), MMLU (+7.1), ARC-C (+7.0).

Sparsity	Method	ARC-C	ARC-E	BoolQ	HSwag	MMLU	OBQA	RTE	Wino	Avg
0%	Full model	.625	.838	.887	.797	.802	.446	.769	.736	.737
25%	EvoESAP	.514	.758	.866	.670	.662	.366	.780	.702	.665
	RCO	.584	.797	.885	.757	.733	.424	.783	.714	.710
50%	EvoESAP	.437	.638	.805	.518	.576	.332	.747	.629	.585
	RCO	.428	.632	.835	.624	.587	.336	.751	.647	.605

At 25% sparsity, RCO improves over EvoESAP on every benchmark. The gains are largest on benchmarks that test world knowledge and reasoning (HellaSwag +8.7, MMLU +7.1, ARC-C +7.0), suggesting that the manifold search better preserves the experts carrying factual and reasoning capacity. At 50% sparsity, EvoESAP retains a small edge on ARC-C (−0.9) and ARC-E (−0.6), while RCO dominates on HellaSwag (+10.6), BoolQ (+3.0), and WinoGrande (+1.8).

C.4 Qwen3-Coder-Next: calibration and allocation ablation

Qwen3-Coder-Next has 512 routed experts per layer with 10 active per token. We evaluate eight RCO variants spanning two calibration domains (coding: evol-codealpaca; general: FineWeb-Edu), two sparsity levels (25%, 50%), and two budget allocation strategies (uniform, nonuniform). All evaluations use vLLM with bf16 and greedy decoding.

Uniform vs. nonuniform allocation. Each variant prunes a fixed fraction of total experts across all layers. Uniform allocation keeps the same number of experts per layer (e.g., 384 at 25%, 256 at 50%). Nonuniform allocation lets the optimizer freely distribute the pruning budget based on calibration loss: critical layers keep more experts, redundant layers are pruned more aggressively. The effect is stark on coding benchmarks (Table 7): at 50% sparsity, nonuniform recovers 97% of HumanEval versus 55% for uniform, a 42-point gap. At 25%, the gap narrows to 8 points (100% vs. 92%). The gap grows with sparsity because uniform allocation forces equal pruning on layers that differ substantially in sensitivity; at low sparsity, even sensitive layers retain enough experts to function.

Coding vs. general calibration. RCO minimizes KL divergence on the calibration dataset, so the choice of calibration data determines which capabilities the pruned model preserves. Coding-calibrated variants (evol-codealpaca) preserve code generation at the cost of general knowledge; general-calibrated variants (FineWeb-Edu) preserve Avg at the cost of coding ability. The trade-off is sharp: general-calibrated variants lose coding ability almost entirely (HumanEval $\leq 6.1\%$ at 25% sparsity), while coding-calibrated variants at 50% sparsity still retain 76% of Avg (Table 8). The asymmetry reflects the base model’s specialization: its coding experts are concentrated in a small subset, easily lost when calibration does not exercise them.

Coding benchmarks. Table 7 reports HumanEval (pass@1) and MBPP (pass@1) with recovery relative to the full model.

Table 7: Qwen3-Coder-Next coding benchmarks across calibration domain and allocation strategy. Recovery is relative to the full (unpruned) model. Bold: best among pruned models per sparsity level. All values in %.

Sparsity	Cal.	Alloc.	HE	rec.	MBPP	rec.
0%	Full model		74.4	–	76.4	–
25%	coding	uniform	68.3	92	68.8	90
	coding	nonunif.	74.4	100	67.8	89
	general	uniform	4.3	6	4.6	6
	general	nonunif.	6.1	8	5.8	8
50%	coding	uniform	40.9	55	53.4	70
	coding	nonunif.	72.0	97	69.0	90
	general	uniform	0.0	0	1.8	2
	general	nonunif.	1.2	2	1.0	1

General benchmarks. Table 8 reports all eight general benchmarks. Avg is the unweighted mean. General-calibrated variants preserve Avg up to 100% at 25% (nonuniform) and 92% at 50% (uniform). At 50% general sparsity, uniform (65.4) slightly outperforms nonuniform (64.4), in contrast to the coding results where nonuniform dominates.

Table 8: Qwen3-Coder-Next general benchmarks (Avg and per-benchmark breakdown). Bold: best among pruned models per sparsity level. All values in %.

Sparsity	Cal.	Alloc.	ARC-C	ARC-E	BoolQ	HSwag	MMLU	OBQA	RTE	Wino	Avg
0%	Full model		60.6	82.1	88.5	77.5	76.7	43.0	76.5	66.6	71.4
25%	coding	uniform	50.1	72.2	86.4	69.0	71.0	38.0	72.9	65.5	65.6
	coding	nonunif.	46.2	66.2	85.1	66.5	68.0	36.2	77.6	64.2	63.8
	general	uniform	60.0	80.7	87.6	78.5	70.4	45.2	75.1	67.7	70.7
	general	nonunif.	61.8	82.2	88.2	77.6	71.2	44.2	76.2	69.9	71.4
50%	coding	uniform	40.3	64.1	78.9	57.8	56.4	35.0	67.1	61.6	57.7
	coding	nonunif.	35.6	55.5	77.6	54.8	54.3	34.0	64.6	60.3	54.6
	general	uniform	54.1	77.1	83.9	70.9	61.0	42.8	67.5	65.8	65.4
	general	nonunif.	52.6	76.2	84.2	70.8	59.5	41.4	67.5	63.5	64.4

D Full MoE Mixed-Precision Quantization Results

Setup. MxMoE calibrates a per-(expert, weight-block, bitwidth) layer-output L2 sensitivity table on FineWeb-Edu (64 samples, 2048 tokens), then solves a per-layer ILP that minimizes the sum of these sensitivities under a memory budget; the gate and up projections share a bitwidth and down is free, and attention is fixed at FP16 by design [Duanmu et al., 2025]. RCO reads the same database and runs the manifold search to minimize calibration KL divergence under the same budget, with one bitwidth decision per expert (all three weights tied), a strictly smaller search space than MxMoE’s per-weight allocation. **RCO*** locks attention at FP16 to match MxMoE’s design exactly, isolating the allocation algorithm; **RCO** additionally treats attention as searchable groups under a total-memory budget. Reproduction details and accounting adjustments are in Appendix F.3.

Aggregate metrics. RCO* outperforms MxMoE on every aggregate metric at every target despite a smaller search space. The Wikitext-2 PPL gap is $-0.43 / -0.18 / -0.04$ at 2.5/3.0/3.5 bits, narrowing toward the FP16 floor (6.79) with more bits, and zero-shot accuracy favors RCO* by 0.3–0.6 points. RCO, which additionally quantizes attention to free expert budget at the same total memory, improves a further 0.01–0.08 Wikitext-2 PPL over RCO* and adds another 0.0–0.4 zero-shot points.

Saturation against uniform allocation. The uniform reference rows in Table 3 bound where mixed precision matters: at 4 bit (uniform), the model is already within 0.7 Wikitext-2 PPL and 0.7 zero-shot points of FP16, so further bit allocation can only yield diminishing returns. RCO at 3.5 bit

matches or beats uniform 4-bit on every metric (W2 6.88 vs 6.91, C4 10.23 vs 10.24, AvgPPL 8.76 vs 8.78, Avg ZS .683 vs .680) while using half a bit less per expert: the manifold has saturated against the post-quantization quality available from this GPTQ database. At 2.5 bit the picture is opposite: uniform 3-bit (W2 7.39, AvgPPL 9.38) is competitive with mixed precision in absolute terms but uses 0.5 more bits per expert, so RCO at 2.5 bit recovers most of uniform 3-bit’s quality at strictly lower memory.

Per-task zero-shot. On individual benchmarks, MxMoE retains an edge on a few (Winogrande at 3.0 and 3.5 bit, ARC-Challenge at 2.5 bit), but the gaps are within 0.4–2.1 points; RCO and RCO* take the remaining benchmarks across all targets, with the largest gains on LAMBADA-OpenAI (+1.4 to +1.6 at 2.5 bit) and ARC-Easy (+2.9 at 3.0 bit).

Table 9: Full Qwen1.5-MoE-A2.7B mixed-precision quantization results, including the 3.0 bit block omitted from Table 3 for space. PPL ($\downarrow$) on Wikitext-2 (W2), C4, FineWeb-Edu (FW); zero-shot accuracy ($\uparrow$, decimals) on six lm-eval-harness tasks; Avg is the unweighted mean.

Bits	Method	Perplexity ($\downarrow$)				Zero-shot accuracy ($\uparrow$)						
		W2	C4	FW	Avg PPL	PIQA	HSwag	ARC-E	ARC-C	Wino	LAMBADA	Avg ZS
FP16	full model	6.79	10.05	9.07	8.64	.805	.773	.690	.445	.695	.713	.687
2.5 bit	MxMoE	7.90	12.44	10.28	10.21	.786	.746	.659	.416	.653	.662	.653
	RCO*	7.47	11.57	9.81	9.61	.793	.750	.652	.411	.651	.677	.656
	RCO	7.40	11.44	9.74	9.53	.792	.752	.667	.409	.665	.676	.660
3.0 bit	uniform	7.39	11.05	9.71	9.38	.789	.761	.650	.436	.663	.689	.664
	MxMoE	7.21	10.90	9.51	9.21	.796	.764	.655	.434	.688	.686	.671
	RCO*	7.03	10.54	9.32	8.96	.799	.766	.684	.442	.673	.700	.677
	RCO	7.02	10.47	9.30	8.93	.798	.767	.677	.453	.667	.700	.677
3.5 bit	MxMoE	6.94	10.35	9.23	8.84	.801	.768	.659	.433	.686	.697	.674
	RCO*	6.90	10.25	9.18	8.77	.799	.770	.687	.445	.673	.705	.680
	RCO	6.88	10.23	9.17	8.76	.801	.771	.686	.453	.684	.706	.683
4 bit	uniform	6.91	10.24	9.20	8.78	.809	.771	.679	.440	.681	.702	.680

E Mixed-Precision Quantization: Ablation Study

We report ablation sweeps for RCO applied to mixed-precision quantization on Qwen3-8B (252 linear layers, 7 bitwidth options: 2, 3, 4, 5, 6, 7, 8) at a target of 2.5 average bits per parameter. All layers are treated as independent groups (no structural grouping unless stated otherwise). Calibration uses FineWeb-Edu with 256 samples at sequence length 2048 unless noted. Evaluation reports perplexity on two held-out corpora: FineWeb-Edu (FW) and C4, each with 131k tokens at sequence length 2048. The FP16 baseline achieves FW=10.96, C4=17.20.

Hardware. All ablation sweeps use an RTX 3090 (24 GB). The baseline comparison (Table 2) reports wall times from an RTX A6000 (48 GB).

E.1 Gumbel samples per step

Each optimization step draws g independent Gumbel-softmax assignments and averages their STE gradients before the optimizer update. More samples reduce gradient variance over the discrete assignment space at the cost of additional forward-backward passes per step.

Table 10: Effect of Gumbel sample count (g). Fixed: $T=200$, $\tau_{\min}=0.01$, $\text{lr}=0.1$, $\text{seed}=42$, $\text{cal}=256$. Hardware: RTX 3090.

g	FW↓	C4↓	Wall	Fwd passes
1	16.46	25.63	23 m	200
4	15.61	24.29	62 m	800
8	15.66	24.13	114 m	1600
16	15.51	24.30	218 m	3200

The largest gain occurs from $g=1$ to $g=4$ (-0.85 FW perplexity). Further increases to $g=16$ yield diminishing and non-monotonic gains on FW (within seed noise), though C4 improves monotonically. Gradient variance over discrete assignments is the primary bottleneck: more samples provide cleaner per-layer signal about which bitwidth to prefer.

E.2 Optimization steps

Temperature anneals exponentially from $\tau_0=1.0$ to $\tau_{\min}$ over T steps, so more steps produce slower annealing.

Table 11: Effect of optimization steps (T). Fixed: $g=4$, $\tau_{\min}=0.01$, $\text{lr}=0.1$, $\text{seed}=42$, $\text{cal}=256$. Hardware: RTX 3090.

T	FW↓	C4↓	Wall	Fwd passes
100	15.62	24.17	35 m	400
200	15.61	24.29	62 m	800
400	15.68	24.34	115 m	1600

Results are non-monotonic: $T=100$ and $T=200$ are nearly identical, while $T=400$ is slightly worse. The optimization converges quickly; additional steps at low temperature accumulate biased STE gradients without improving the solution.

E.3 Samples vs. steps at fixed compute

For a fixed compute budget (same number of forward passes), is it better to increase Gumbel samples per step or to run more steps?

Table 12: Compute-matched configurations. $g=32$, $T=50$ achieves the same quality as $g=16$, $T=200$ at half the compute (1600 vs. 3200 forward passes). Note: $g=32$ uses $\text{seed}=2$; others use $\text{seed}=42$. Hardware: RTX 3090.

g	T	FW↓	C4↓	Fwd passes	Wall
4	200	15.61	24.29	800	62 m
8	100	15.68	24.28	800	61 m
16	200	15.53	24.13	3200	202 m
32	50	15.56	24.10	1600	110 m

Sample efficiency dominates step count. The STE gradient benefits more from averaging over diverse discrete assignments than from additional optimization iterations at low temperature.

E.4 Temperature minimum

Temperature decays exponentially from $\tau_0=1.0$ to $\tau_{\min}$. Lower $\tau_{\min}$ forces sharper assignments earlier; higher $\tau_{\min}$ leaves more uncertainty at convergence. The final discrete assignment is extracted via budget-constrained argmax regardless of $\tau_{\min}$.

Table 13: Effect of $\tau_{\min}$. Fixed: $g=4$, $T=200$, $\text{lr}=0.1$, $\text{seed}=42$, $\text{cal}=256$. Hardware: RTX 3090.

$\tau_{\min}$	FW↓	C4↓
0.05	15.75	24.55
0.01	15.61	24.29
0.005	15.77	24.42
0.001	16.26	25.20

$\tau_{\min}=0.01$ is optimal. Higher values leave too many layers with soft, undecided assignments. Lower values force premature hard decisions before layers have accumulated sufficient gradient signal, locking in suboptimal assignments early.

E.5 Learning rate

The Adam optimizer step size in logit space.

Table 14: Effect of learning rate. Fixed: $g=4$, $T=200$, $\tau_{\min}=0.01$, $\text{seed}=42$, $\text{cal}=256$. Hardware: RTX 3090.

LR	FW↓	C4↓
0.05	15.66	24.33
0.1	15.61	24.29
0.2	15.63	24.40

Results are robust across a $4\times$ range (within seed noise of ~ 0.2 PPL). The manifold projection removes the budget-changing gradient component, and the retraction prevents constraint violation, leaving the optimizer in a well-conditioned subspace.

E.6 Antithetic sampling

Antithetic sampling pairs each uniform draw U with $1-U$, producing two negatively correlated Gumbel variates per draw. This halves the number of independent assignments explored per step.

Table 15: Effect of antithetic sampling. Fixed: $g=4$, $T=200$, $\tau_{\min}=0.01$, $\text{lr}=0.1$, $\text{seed}=42$, $\text{cal}=256$. Hardware: RTX 3090.

Antithetic	FW↓	C4↓
No	15.61	24.29
Yes	15.97	24.63

Antithetic sampling hurts. With $g=4$, the two base draws plus two mirrored draws reduce diversity in the discrete assignment space. The variance reduction from negative correlation does not compensate for the reduced exploration. Independent samples explore the combinatorial space more effectively.

E.7 Seed variance

The random seed affects calibration data ordering, Gumbel noise sequences, and optimizer initialization.

Table 16: Seed variance. Fixed: $g=4$, $T=200$, $\tau_{\min}=0.01$, $\text{lr}=0.1$, $\text{cal}=256$. Hardware: RTX 3090.

Seed	FW↓	C4↓
2	15.57	24.24
42	15.61	24.29
0	15.83	24.84
1	15.96	25.23
3	15.98	24.71

FW perplexity: mean=15.79, std=0.18, range=0.41. The range exceeds most hyperparameter effects except Gumbel sample count. Approximately 180 of 252 layers converge to the same assignment across seeds; the remaining ~ 70 ambiguous layers (where multiple bitwidths yield similar KL) account for the variance. More Gumbel samples reduce this variance by providing better gradient signal for ambiguous layers.

E.8 Layer grouping

Structural grouping forces layers with similar roles (e.g., gate_proj and up_proj, or q/k/v_proj) to share a bitwidth, reducing the search space from 252 to ~ 144 groups.

Table 17: Effect of layer grouping. Fixed: $g=4$, $T=200$, $\tau_{\min}=0.01$, $\text{lr}=0.1$, $\text{seed}=42$. Hardware: RTX 3090.

Grouping	Cal	FW↓	C4↓
None	256	15.60	24.14
Grouped	256	16.11	25.36
None	512	15.82	24.83
Grouped	512	16.94	26.98
None	1024	16.14	25.13
Grouped	1024	16.40	25.73

No grouping consistently outperforms grouped assignment by 0.3–1.1 FW perplexity. Individual layers within structural groups can have different quantization sensitivity, so forcing a shared bitwidth is suboptimal. RCO scales well to the full 252-dimensional search space without grouping.

E.9 Calibration samples

More calibration sequences produce a less noisy KL objective but require proportionally more RAM for cached reference log-probabilities (~ 159 GB for 256 samples).

Table 18: Effect of calibration set size. Note: $\text{cal}=512$ runs use $\text{lr}=0.2$; comparison is not perfectly controlled. Hardware: RTX 3090.

Cal	Config	FW↓	C4↓
256	$g=8$, $T=200$, $\text{lr}=0.1$	15.78	24.48
512	$g=8$, $T=200$, $\text{lr}=0.2$	15.65	24.51
256	$g=4$, $T=200$, $\text{lr}=0.1$	15.81	24.73
512	$g=4$, $T=200$, $\text{lr}=0.2$	16.23	25.38

No clear benefit from doubling calibration data. The dominant source of variance is the discrete assignment landscape (seed variance, Section E.7), not calibration noise. At 256 samples (524k tokens), the gradient signal is sufficient for the ~ 252 decisions.

E.10 Scaling: Gumbel samples with fixed steps

Table 19 summarizes the compute-quality tradeoff as Gumbel sample count increases, alongside EvoPress.

Table 19: Scaling behavior. RCO with $g=32$, $T=50$ matches or exceeds EvoPress on both FW and C4. The compute-optimal strategy maximizes samples per step while minimizing step count. RCO rows on RTX 3090; EvoPress row from the matched-protocol reproduction in Table 2 (RTX A6000).

Method	g	T	FW↓	C4↓	Wall
RCO	1	200	16.46	25.63	23 m
RCO	4	200	15.61	24.29	62 m
RCO	8	200	15.66	24.13	114 m
RCO	16	200	15.53	24.30	218 m
RCO	32	50	15.56	24.10	110 m
EvoPress	=	100 gen	15.64	24.63	11–14 h

E.11 Recommended defaults

Table 20 summarizes the recommended hyperparameters based on the ablations above.

Table 20: Recommended RCO hyperparameters for mixed-precision quantization.

Parameter	Recommended	Robust range	Note
Gumbel samples (g)	16–32	4–64	More is better; diminishing returns
Steps (T)	50–200	50–400	Non-monotonic; fewer is fine with more samples
$\tau_{\min}$	0.01	0.005–0.05	Too low: premature decisions
Learning rate	0.1	0.05–0.2	Robust due to manifold projection
Antithetic	No	-	Reduces assignment diversity
Calibration	256	128–512	Not the bottleneck
Layer grouping	None	-	Layers differ in sensitivity

F Baseline Reproduction Details

All evolutionary search baselines reported in this paper are our own reproductions under a matched evaluation protocol, not numbers taken from the original papers. This section documents the reproduction setup for each.

F.1 EvoESAP

All EvoESAP numbers are produced using the authors’ released code¹ under our evaluation pipeline: same models, same FineWeb-Edu calibration data and KL-divergence fitness used for RCO, same lm-eval-harness configuration for downstream benchmarks, same compute environment. We use REAP as the within-layer importance criterion since it is the strongest of the four criteria evaluated by Liu et al. [2026]. We run EvoESAP at the generation counts recommended in their paper (50 for OLMoE-1B-7B, 10 for Qwen3-30B-A3B). Our EvoESAP numbers therefore reflect performance under a matched protocol with RCO and may differ from the numbers in Liu et al. [2026], who calibrate REAP on evol-codealpaca-v1 and search on tulu-3-sft-personas-math; the difference reflects calibration data choice, not algorithmic performance.

F.2 EvoPress

All EvoPress numbers for Qwen3-8B are our own reproductions, run using the authors’ released code² under our evaluation pipeline (FineWeb-Edu calibration, KL divergence fitness, 256 calibration samples, same eval setup as RCO). Sieberling et al. [2025] do not include Qwen3-8B in their reported experiments. We run EvoPress for 100 generations, matching the search budget recommended by Sieberling et al. [2025]. RCO and EvoPress are measured in the same compute environment.

¹<https://github.com/ZongfangLiu/EvoESAP>

²<https://github.com/IST-DASLab/EvoPress>

F.3 MxMoE

All MxMoE numbers for Qwen1.5-MoE-A2.7B are produced from the authors’ released code³ under our evaluation pipeline: a shared GPTQ weight database (bitwidths 2–8, group size 128, asymmetric, no Hadamard) read by both methods, MxMoE’s per-bitwidth sensitivity calibration command unchanged (`mxmoe.quant.quant calib -metric layer_out_norm`, FineWeb-Edu, 64 samples $\times$ 2048 tokens), the per-layer ILP allocator (binary $x_{e,n,s}$, objective $\min \sum_{e,n,s} x_{e,n,s} \cdot \delta_{e,n,s}$ on the layer-output sensitivity δ , with `gate=up` coupling and down free; attention is left at FP16 since their pipeline hardcodes `w_bits=16` for it), the Wikitext-2 PPL protocol (100 chunks of 4096 tokens), the zero-shot task list (PIQA, HellaSwag, ARC-Easy, ARC-Challenge, Winogrande, LAMBADA-OpenAI), and our harness for both PPL and zero-shot. Our only modification is to remove the +0.25-bit group-128 metadata overhead from the bit accounting so that both methods report the same bit rate. We omit three MxMoE features that are orthogonal to the allocation algorithm: the online Hadamard transform (`gptq-had`, their flagship variant; the weight database is our own and would need to be rebuilt with Hadamard, though Hadamard would benefit RCO equally), joint accuracy and throughput optimization (does not affect the accuracy results we report), and weight-and-activation quantization (W4A4, W8A8; our comparison is weight-only).

G Full MCKP Results

We evaluate on 13 scenarios (3 random instances per scenario, 5000 steps, learning rate 0.01) spanning correlated costs, tight budgets, adversarial instances, under-budget optima, and a large-scale instance ($N=1000$, $K=32$). Table 21 reports the five scenarios on which the methods diverge; the remaining eight place all four methods (equality, slack, Lagrangian, augmented Lagrangian) within instance-to-instance noise of each other. Gap denotes the percentage below the DP optimum for the greedy-repaired discrete assignment; violation is the mean absolute constraint violation $|C(\alpha) - B|$ averaged over all optimization steps.

Table 21: MCKP benchmark results on the scenarios that distinguish methods (5000 steps, 3 random instances per scenario). On the remaining eight scenarios all methods land within instance noise of each other. Gap: % below DP optimum on the greedy-repaired discrete assignment (mean $\pm$ std over instances). Violation: mean $|E[\text{cost}] - B|$ averaged over all steps. Budget manifold maintains constraint violation $< 10^{-8}$ throughout. [†]The slack variant of Budget manifold recovers the DP optimum on both *cheap optimal* and *mixed slack* (gap $< 0.01\%$), while equality-constrained methods waste the surplus budget on inferior options.

Scenario	Budget manifold		Lagrangian		Aug. Lagrangian	
	Gap%	v	Gap%	v	Gap%	v
<i>Scale: Lagrangian penalty degrades with N, K.</i>						
<i>large</i> ($N=200$, $K=16$)	0.02 $\pm$ 0.02	$< 10^{-8}$	0.02 $\pm$ 0.02	0.08	0.53 $\pm$ 0.12	0.03
<i>float costs</i> ($N=200$, $K=16$)	0.02 $\pm$ 0.02	$< 10^{-8}$	0.02 $\pm$ 0.02	0.08	0.64 $\pm$ 0.04	0.03
<i>huge</i> ($N=1000$, $K=32$)	0.00 $\pm$ 0.00	$< 10^{-8}$	37.56 $\pm$ 1.78	1.15	4.28 $\pm$ 0.20	0.12
<i>Under-budget optima: equality methods waste budget; the slack variant[†] recovers the optimum.</i>						
<i>cheap optimal</i>	48.71 $\pm$ 1.41	$< 10^{-8}$	48.75 $\pm$ 1.14	0.09	48.75 $\pm$ 1.14	0.01
<i>mixed slack</i>	12.57 $\pm$ 1.66	$< 10^{-8}$	12.35 $\pm$ 1.65	0.08	12.35 $\pm$ 1.48	0.00

The results fall into two regimes that distinguish methods, plus a residual.

Scale (*large*, *float costs*, *huge*). Lagrangian penalty performance degrades sharply with N and K : at $N=1000$ ($K=32$), Lagrangian plateaus at a 37.6% gap while Budget manifold reaches the DP optimum exactly. Augmented Lagrangian is more robust at scale but still trails Budget manifold by an order of magnitude on the largest instance and by 25–30 $\times$ on the moderate-scale ones. Across all three scenarios, Budget manifold’s repaired gap stays at $\leq 0.02\%$.

³<https://github.com/cat538/MxMoE>

Under-budget optima (*cheap optimal, mixed slack*). When the value-maximizing assignment costs less than the budget, equality-constrained methods are forced to spend the remaining budget on inferior options, losing up to half the optimal value. The slack variable resolves this completely: Budget manifold (inequality) recovers the DP optimum exactly. This scenario is relevant in practice when budget targets are set conservatively.


Easy / tied scenarios. On *small easy, medium, tight budget, adversarial, and nonuniform*, the four methods reach comparable final gaps and comparable compute (within instance-to-instance noise of each other). Constraint feasibility nonetheless separates them throughout training: Budget manifold (eq.) maintains $|C(\alpha) - B| < 10^{-8}$ in every scenario, whereas Lagrangian methods carry a mean violation of order 10^{-1} (Figures 6, 7).

Compute-matched comparison across scenarios. Final-step gap masks substantial differences in convergence speed. Table 22 reports the number of steps each method needs to reach within 1% of the DP optimum on the eight scenarios that distinguish them. Three regimes emerge. At scale ($N \geq 200$), augmented Lagrangian does not converge to 1% within the 5000-step budget on any of *large, float costs*, or *huge*; vanilla Lagrangian additionally plateaus indefinitely on *huge*. On hard small-scale instances (*boundary, correlated, correlated tight*), the manifold reaches the threshold $2\text{--}5\times$ faster than augmented Lagrangian and $1.5\text{--}2.5\times$ faster than vanilla Lagrangian. On the under-budget scenarios (*cheap optimal, mixed slack*), only the slack variant converges; the equality manifold and both Lagrangian baselines plateau indefinitely on the under-budget portion of the surface.


Table 22: Compute (in optimization steps) for each method to reach within 1% of the DP optimum, on the eight scenarios that distinguish methods (averaged over 3 random instances per scenario; *never* indicates the threshold is not reached within 5000 steps). Bold marks the fastest method in each row. [†]Under-budget scenarios: equality methods structurally cannot leave the budget surface; the reported value is the slack variant, which converges in ~ 500 steps.

Scenario	Manifold	Lagrangian	Aug. Lagrangian
<i>Scale: $N \geq 200$.</i>			
<i>huge</i> ($N=1000, K=32$)	594	<i>never</i>	<i>never</i>
<i>large</i> ($N=200, K=16$)	583	583	<i>never</i>
<i>float costs</i> ($N=200, K=16$)	583	585	<i>never</i>
<i>Hard small-scale: correlated values or boundary-clustered structure.</i>			
<i>boundary</i> ($N=50, K=8$)	185	453	900
<i>correlated</i> ($N=50, K=8$)	381	581	1444
<i>correlated tight</i> ($N=50, K=8$)	308	482	1029
<i>Under-budget optima: only the slack variant of the manifold converges.</i>			
<i>cheap optimal</i> ($N=50, K=8$)	562[†]	<i>never</i>	<i>never</i>
<i>mixed slack</i> ($N=50, K=8$)	514[†]	<i>never</i>	<i>never</i>

Figures 6 and 7 show per-instance convergence traces for representative scenarios. The raw constraint violation traces (panel c in each figure) make the oscillation of Lagrangian methods visible: the penalty weight overshoots and undershoots the budget repeatedly, while Budget manifold maintains violation at $\sim 10^{-8}$ throughout.




(a) Adversarial scenario ($N=50, K=8$, single instance). Costs clustered near budget/ N , values nearly tied. The Lagrangian penalty oscillates persistently (panel c), never satisfying the constraint. Budget manifold converges to the DP optimum with zero violation.




(b) Correlated tight scenario ($N=50, K=8$, single instance). Strong cost-value correlation with budget only 15% above minimum. Budget manifold converges faster (panels a,b) and maintains 10^{-8} violation (panel d) versus 10^{-2} for Lagrangian methods.

Figure 6: MCKP convergence traces for adversarial and correlated tight scenarios.



(a) Boundary-clustered scenario ($N=50, K=8$, single instance). Many near-optimal solutions packed at the budget boundary. Log-scale convergence (panel b) shows Budget manifold reaching 10^{-6} gap while Lagrangian methods plateau at 10^{-3} .



(b) Cheap-optimal scenario ($N=50, K=8$, single instance) with slack variable. The DP optimum costs well under budget. Equality-constrained methods (Budget manifold (eq.), Lagrangian, Aug. Lagrangian) all converge to $\sim 50\%$ of optimal (panel a). Budget manifold with slack finds the true optimum, with the slack variable absorbing the unused budget (panel c: $C(\alpha) - B \rightarrow -276$).

Figure 7: MCKP convergence traces for boundary-clustered and cheap-optimal scenarios.